


Raw kept the strongest floor; the manifest halved context

1

Corrected isolated Bloomberg continuation experiment · five runs per condition · 21 hidden contracts



WHAT THE CORRECTED RUNS SHOWED

52.9%

less context

Report + manifest removed 28,569 prompt tokens while retaining a 16.0/21 mean.

16.8 · 16.0

raw · manifest means

Raw kept the strongest floor: 16. Manifest ranged from 14 to 20.

34.0%

Shake reduction


Shake kept all 120 messages and ranged from 12 to 20 contracts.

Decision: raw is the safety default; report + manifest is the explicit context-pressure option.

Report + manifest bought the most context per unit of work

2

Corrected five-replicate workload and cost comparison



Report + manifest used the fewest mean tools, finished fastest, and cost least in these runs.

Scoped Shake reduced prompt size, but required the most work and had the widest correctness spread.

PRODUCT DECISION

SAFETY DEFAULT

Keep raw history

Highest mean (16.8), strongest floor (16).
Preserve load-bearing trajectory detail.

CONTEXT PRESSURE

Report + manifest

52.9% smaller prompt, 16.0 mean, 14 minimum, and the best observed workload and cost profile.

SPECIAL CASE

Scoped Shake

Use when chronology must stay visible. It saved 34.0%, but did not beat the semantic candidate here.

QUALIFICATIONS

- All 20 valid runs passed the offline suite, Ruff, basedpyright, and ty in the external evaluator.
- Every run missed the same server-not-tested/schema-hash contract; relative differences remain visible, but absolute scores are capped.
- Five runs per arm and one task archetype support a directional decision—not automatic default promotion.

• One cwd-mismatch setup pilot and one provider-timeout attempt were rejected, retained separately, and excluded before aggregation.
Source: checkpoint-seal-bloomberg-clean-v2-0.144.1-medium/analysis/aggregate.json · 20 valid runs.